

Impletementing Ridge and Lasso Regression

- 1. Importing Libraries Necessary for the project.
- 2. Defining graph function to be lter used to visualize Data
- 3. Understanding Data, it's Coloumns and type.
- 4. Cleaning Data for use and Data preprocessing, by converting object type to int.
- 5. Diving Data into Training and Target Data Sets.
- 6. Visualizing Training Data Set using Histograms, Heatmaps and displots.
- 7. Removing neglectable Features that we know after Visualization.
- 8. Visualizing Again to be sure that the data is now ready to be used.
- 9. Splitting Training and Testing Data From Training Data.
- 10. Training our model using Ridge Regression and then calculating it's accuracy.
- 11. Visualizing our Coeffecients and effect of alpha on coefficients.
- 12. Applying trianed model on Testing Data and predicting target variable.
- 13. Training another model using Lasso Regression and calculating it's accuracy.
- 14. Visualizing our Coefficients and effect of alpha on Coefficients.
- 15. Applying trained model on Testing Data and predicting target variable.
- 16. Comparing and visualizing outputs from Ridge and Lasso Regression.
- 17. And lastly, we plot the difference between lasso and ridge regression.

Importing Libraries

```
In [1]:  
  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.linear_model import Ridge  
from sklearn.linear_model import Lasso  
from sklearn.metrics import mean_squared_error, r2_score
```

Defining Functions

1. Function for Heatmap

```
In [2]:  
  
def Heat(Datax):  
    correlate = Datax.corr()  
    fig, ax = plt.subplots(figsize = (20,10))  
    ax = sns.heatmap(corelate, annot = True)  
    plt.show()
```

1. Function for distplots

```
In [3]:  
  
def dis(Dcol):  
    for i in Dcol.columns:  
        fig, ax = plt.subplots(figsize = (6, 3))  
        xlabel = 'a'  
        ax = sns.distplot(Dcol.loc[:,i])  
        plt.show()
```

1. Function For Bar Graphs

```
In [4]:  
  
def Bar(x,y):  
    fig, ax = plt.subplots(figsize=(18,10))  
    ax = plt.bar(x, y)  
    plt.style.use('ggplot')  
    plt.show()
```

Reading Data

```
In [5]:  
  
Data = pd.read_excel('Hitters.xlsx')
```

Type of Data for each column

```
In [6]:  
  
Data.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 322 entries, 0 to 321  
Data columns (total 20 columns):  
AtBat      322 non-null int64  
Hits       322 non-null int64  
HmRun      322 non-null int64  
Runs       322 non-null int64  
RBI        322 non-null int64
```

Walks 322 non-null int64
Years 322 non-null int64
CAtBat 322 non-null int64
CHits 322 non-null int64
CHmRun 322 non-null int64
CRuns 322 non-null int64
CRBI 322 non-null int64
CWalks 322 non-null int64
League 322 non-null object
Division 322 non-null object
PutOuts 322 non-null int64
Assists 322 non-null int64
Errors 322 non-null int64
Salary 263 non-null float64
NewLeague 322 non-null object
dtypes: float64(1), int64(16), object(3)
memory usage: 50.4+ KB

Converting type object into int64 for use.

In [7]:

```
for i in range(len(Data)):
    if (Data.loc[i, 'League'] == 'A'):
        Data.loc[i, 'League'] = 50
    else:
        Data.loc[i, 'League'] = 60

    if (Data.loc[i, 'NewLeague'] == 'A'):
        Data.loc[i, 'NewLeague'] = 50
    else:
        Data.loc[i, 'NewLeague'] = 60

    if (Data.loc[i, 'Division'] == 'E'):
        Data.loc[i, 'Division'] = 50
    else:
        Data.loc[i, 'Division'] = 60
```

Extracting Training Data From Data Set

In [8]:

```
Data_train = pd.DataFrame(Data)
```

In [9]:

```
for i in range(0, len(Data_train)):
    if (str(Data_train.loc[i, 'Salary']) == 'nan'):
        Data_train.drop([i], axis = 0, inplace = True)
```

Extracting Target Data From Data Set

In [10]:

```
Data_target = pd.DataFrame(Data)
```

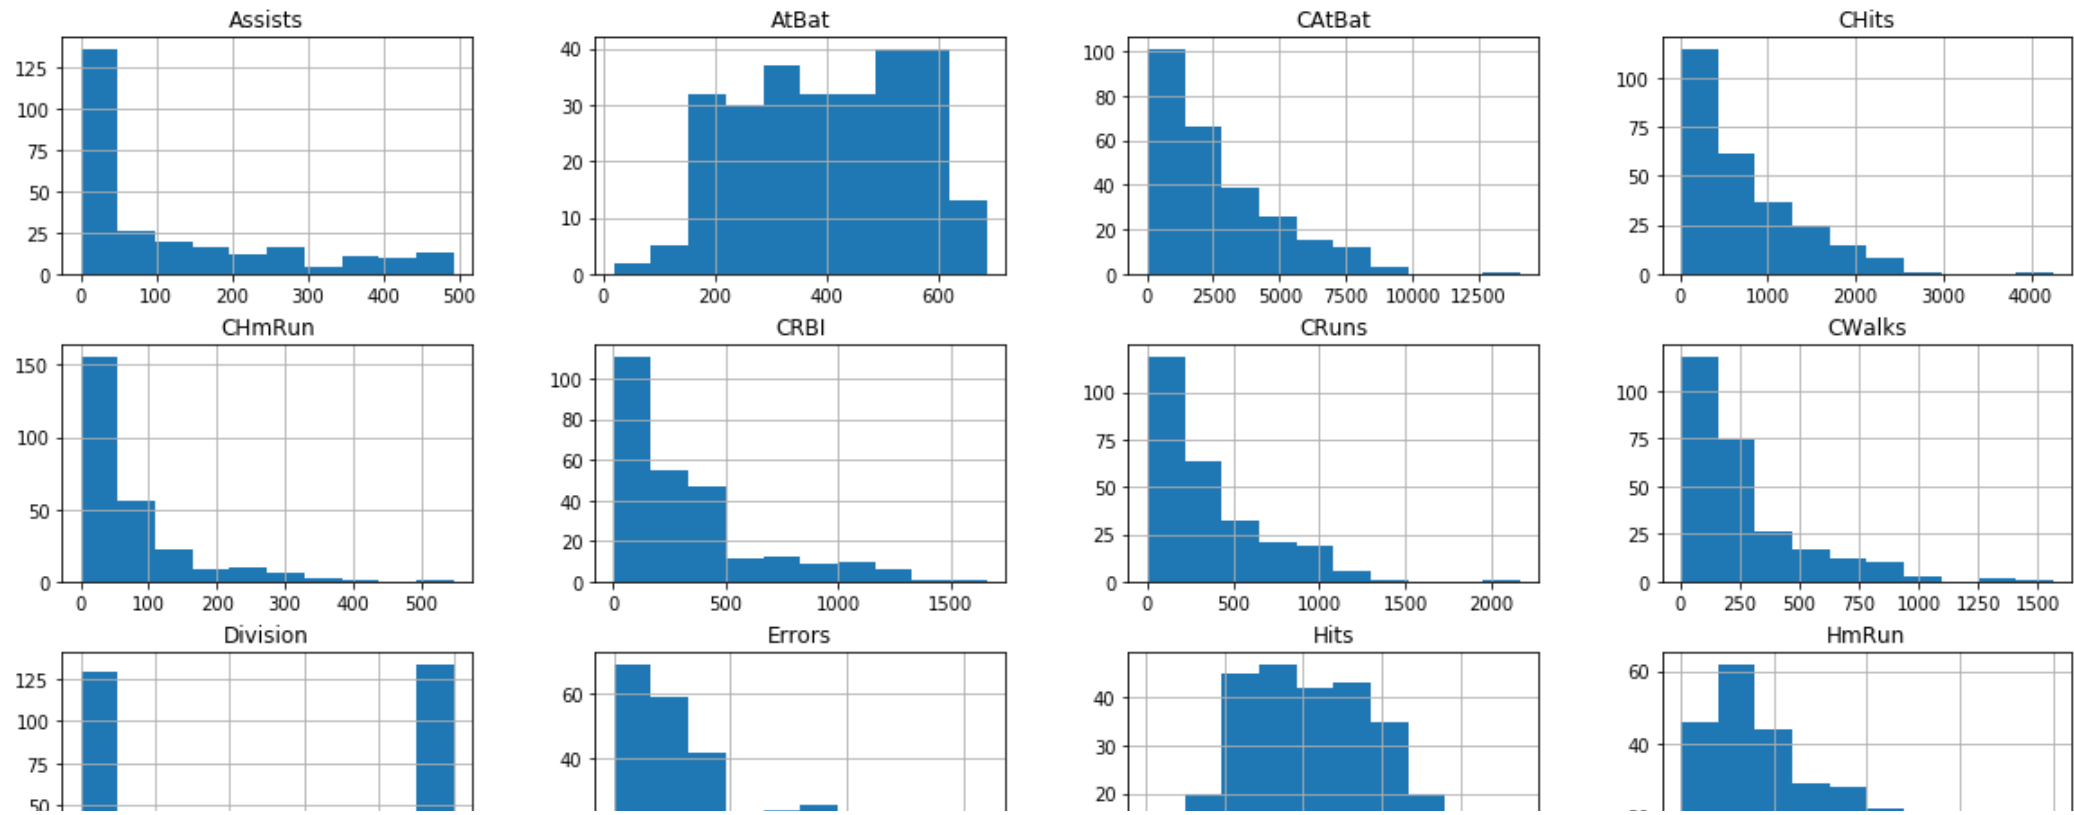
In [11]:

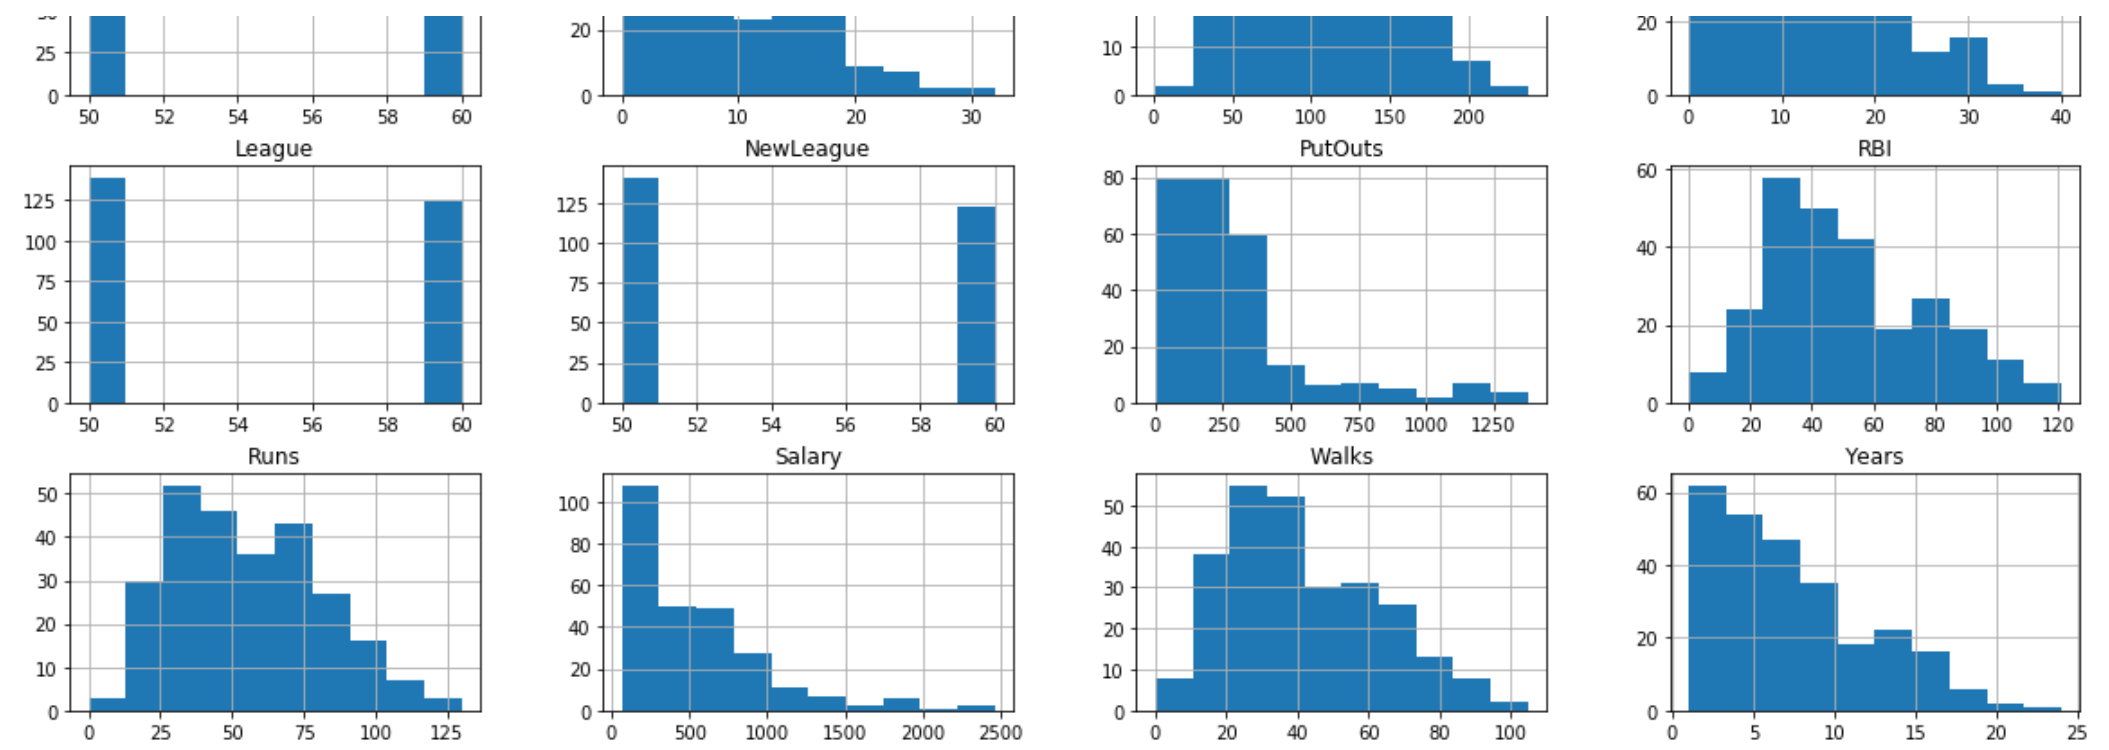
```
for i in range(0, len(Data_target)):
    if (str(Data_target.loc[i, 'Salary']) != 'nan'):
        Data_target.drop([i], inplace = True)
```

Visualizing Data with Histogram for each Columns

In [12]:

```
Data_train.hist(figsize = (20,15))
plt.show()
```






Checking for least required Features with Heatmap

In [13]:

```
Heat(Data_train)
#dis(Data_train)
```

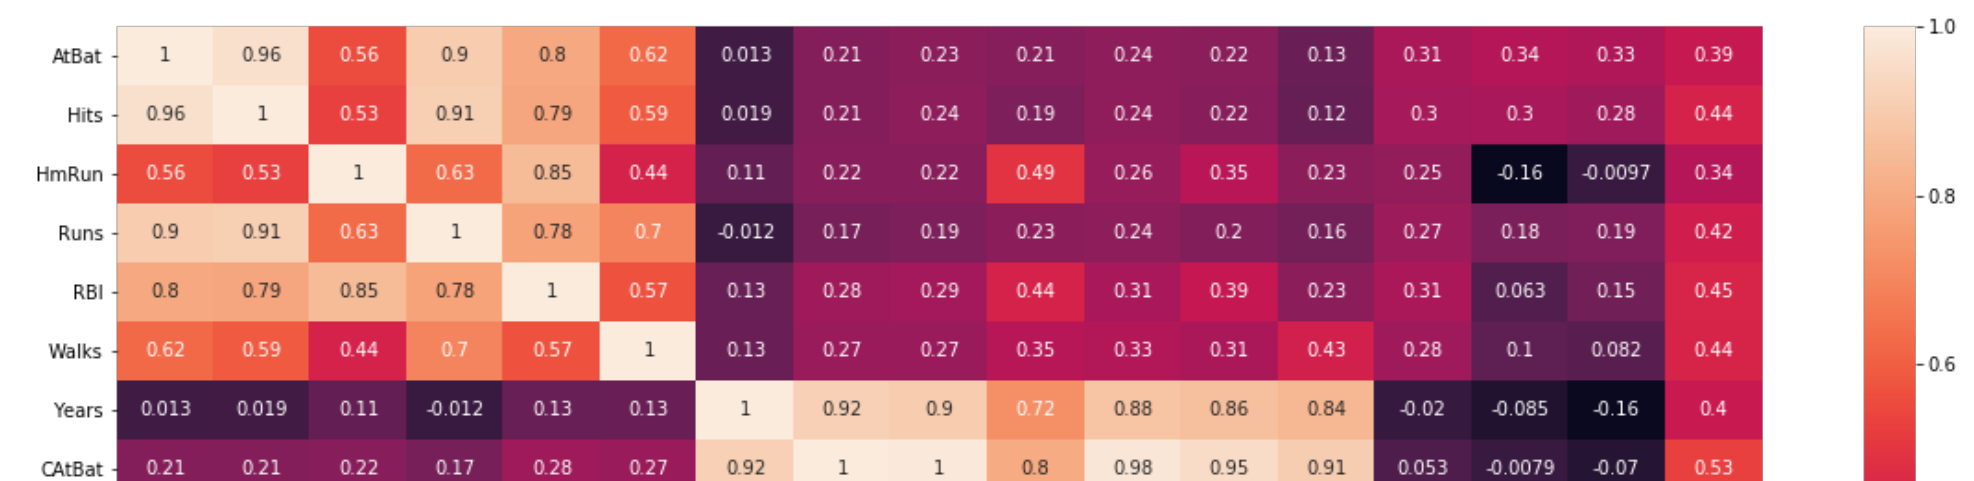


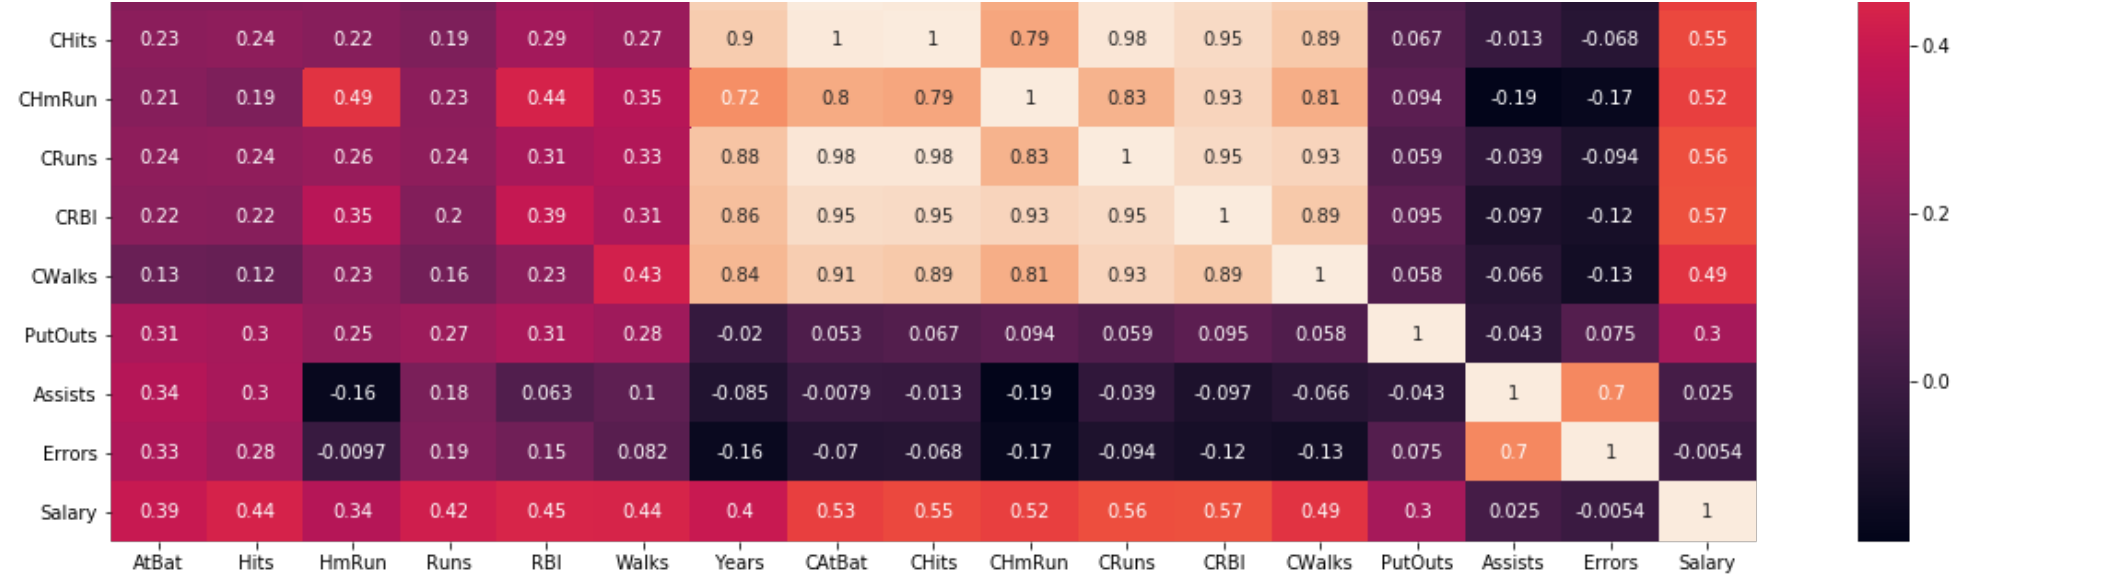
Dropping Features that our bad

In [14]:

```
Data_train = Data_train.drop(['Division'],axis = 1)
Data_train = Data_train.drop(['League'],axis = 1)
Data_train = Data_train.drop(['NewLeague'],axis = 1)
Data_train.head()
print('After removing bad features from our data.')
Heat(Data_train)
#dis(Data_train)
```

After removing bad features from our data.





In [15]:

```
Data_target = Data_target.drop(['Division'],axis = 1)
Data_target = Data_target.drop(['League'],axis = 1)
Data_target = Data_target.drop(['NewLeague'],axis = 1)
Target_array = np.array(Data_target)
X_target = Target_array[:, :-1]
```

Splitting Training and Test Data within Training Data

In [16]:

```
Data_array = np.array(Data_train)

DFX = pd.DataFrame(Data_train.loc[:, :,'Salary'])
DFX = DFX.drop(['Salary'],axis =1)

X_train = Data_array[:235, :-1]
X_train_test = Data_array[235:, :-1]

Y_train = Data_array[:235, -1]
Y_train_test = Data_array[235:, -1]
```

Ridge Regression

1. Training model 'ridge'

In [17]:

```
ridge = Ridge(alpha = .0008,normalize = True)
ridge.fit(X_train,Y_train)
```

Out[17]:

Ridge(alpha=0.0008, copy_X=True, fit_intercept=True, max_iter=None, normalize=True, random_state=None, solver='auto', tol=0.001)


1. Visualizing Coefficients

In [18]:

```
ridge_coef = pd.DataFrame()
ridge_coef["Columns"] = DFX.columns
ridge_coef["Coefficient"] = pd.Series(ridge.coef_)
print(ridge_coef)
Bar(ridge_coef["Columns"], ridge_coef["Coefficient"])
```

	Columns	Coefficient
0	AtBat	-2.142672
1	Hits	6.811597
2	HmRun	0.451745
3	Runs	0.150867
4	RBI	0.492329
5	Walks	5.001096
6	Years	-2.410828
7	CAtBat	-0.157967
8	CHits	0.025681
9	CHmRun	-0.450001
10	CRuns	1.334692
11	CRBI	1.018829
12	CWalks	-0.676457
13	PutOuts	0.326991
14	Assists	0.371144
15	Errors	-2.421001






Calculating Rmse for Training and Testing Data and Plotting original and predicted values For Ridge Regression.

In [19]:

```
Y_pred_train = ridge.predict(X_train)
rmse_train = np.sqrt(mean_squared_error(Y_train, Y_pred_train))
Y_pred_train_test = ridge.predict(X_train_test)
rmse_train_test = np.sqrt(mean_squared_error(Y_train_test, Y_pred_train_test))
print('For Training Data', rmse_train)
print('For Testing Data', rmse_train_test)
plt.scatter(Y_train, Y_pred_train,label = 'Training')
plt.legend(['Training'],loc = 'upper right')
plt.plot(Y_train_test, Y_pred_train_test, label = 'Testing', color = 'blue')
plt.xlabel('Value of original Data Set')
plt.ylabel('Predicted value')
plt.legend(loc = 'upper right')
plt.show()
```

For Training Data 311.99234450323274
For Testing Data 297.5623689562631



Predicting Values for Target Data Set

In [20]:

```
Y_Target_R = ridge.predict(X_target)
print('Predicted values with Ridge Regression\n',Y_Target_R)
```

Predicted values with Ridge Regression
[137.41726007 79.24885396 831.28208311 393.31817108 672.97758097
1258.15056534 198.62930405 761.92286909 187.35827331 352.71136491
342.69408879 291.78501091 1303.69967368 106.29679134 419.11112635
112.06134179 138.28995929 445.53136008 293.75582351 269.72532734
859.6532685 124.49836226 321.68775774 676.86783454 369.39267306
531.31579572 1106.88301395 57.50193368 157.93717162 974.07141216
523.19368262 167.1385725 378.83056175 243.13449124 175.8325892
592.37909595 155.47163531 218.18439234 244.0508993 305.342201
118.38249318 659.73830963 172.16881848 342.15223163 349.66127328
416.72936717 336.8657662 1069.49869477 597.18922427 165.2371818
118.9847288 396.57736752 233.41172948 138.6196837 675.12401357
139.32521769 1343.85373836 961.39803729 240.98228541]

Lasso Regression

1. Training Model

In [21]:

```
lasso = Lasso(alpha = 0.01, normalize = True)
lasso.fit(X_train, Y_train)
```

C:\Users\nabhr\Anaconda3\lib\site-packages\sklearn\linear_model\coordinate_descent.py:475: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 434665.08528796956, tolerance: 4900.630711327813 positive)

Out[21]:


```
Lasso(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=1000,
      normalize=True, positive=False, precompute=False, random_state=None,
      selection='cyclic', tol=0.0001, warm_start=False)
```

1. Visualizing Coefficients

In [22]:

```
lasso_coef = pd.DataFrame()
lasso_coef["Columns"] = DFX.columns
lasso_coef["Coefficient"] = pd.Series(lasso.coef_)
print(lasso_coef)
Bar(lasso_coef["Columns"], lasso_coef["Coefficient"])
```

	Columns	Coefficient
0	AtBat	-2.136390
1	Hits	6.892861
2	HmRun	0.756205
3	Runs	-0.000000
4	RBI	0.351007
5	Walks	5.013487
6	Years	-1.312714
7	CAtBat	-0.169843
8	CHits	-0.000000
9	CHmRun	-0.608980
10	CRuns	1.423631
11	CRBI	1.094356
12	CWalks	-0.687808
13	PutOuts	0.326581
14	Assists	0.371738
15	Errors	-2.294680





Calculating Rmse for Training and Testing Data and Plotting original and predicted values For Ridge Regression.

In [23]:

```
Y_pred_train = lasso.predict(X_train)
rmse_train = np.sqrt(mean_squared_error(Y_train, Y_pred_train))
Y_pred_train_test = lasso.predict(X_train_test)
rmse_train_test = np.sqrt(mean_squared_error(Y_train_test, Y_pred_train_test))
print('For Training Data', rmse_train)
print('For Testing Data', rmse_train_test)
plt.scatter(Y_train, Y_pred_train,label = 'Training')
plt.legend(['Training'],loc = 'upper right')
plt.plot(Y_train_test, Y_pred_train_test, label = 'Testing', color = 'blue')
plt.xlabel('Value of original Data Set')
plt.ylabel('Predicted value')
plt.legend(loc = 'upper right')
plt.show()
```

For Training Data 311.92758702361226
For Testing Data 297.8098719558876





Predicting Values for Target Data set using Lasso

In [24]:

```
Y_Target_L = lasso.predict(X_target)
print('Predicted Values with Lasso Regression\n',Y_Target_L)
```

Predicted Values with Lasso Regression

[137.48651377	79.83018758	825.89860059	399.04197483	672.51948672
1259.216157	202.55689726	759.91072103	187.24527981	340.51148014	
346.69760972	288.44468137	1310.52350329	108.68410859	418.61689699	
113.75344441	137.63928474	444.81789638	294.82235331	270.16923181	
851.19491314	123.53384939	323.34734763	678.67670696	369.20363602	
531.8974074	1115.38423638	59.95248731	157.34874694	979.12623919	
522.72115371	163.53223762	377.94591177	249.01037284	175.77863305	
592.87498155	160.32328775	219.16873039	244.09099604	303.8461139	
117.73107216	670.13089692	172.86506622	339.72185822	346.75516828	
415.71839539	338.09077985	1065.45512225	595.32878217	164.71266516	
118.00283589	398.4278276	233.5796233	139.67747014	677.71112028	
139.31105202	1349.01551764	967.68376774	244.17159622]		

Comparing Predicted Values of Target Datasets by Ridge and Lasso Regression.

In [25]:

```
A = []
for i in range(len(Y_Target_L)):
    A.append(abs(Y_Target_L[i] - Y_Target_R[i]))
max(A)
```

Out[25]:


12.199884765364345


Comparing Ridge and Lasso Coefficients

In [26]:

```
print('Difference between coefficients from ridge and lasso regression')
Bar(ridge_coef['Columns'], lasso_coef['Coefficient'] - ridge_coef['Coefficient'])
X = np.arange(16)
fig = plt.figure(figsize=(18,10))
ax = fig.add_axes([0,0,1,1])
ax.bar(X + 0.0, lasso_coef['Coefficient'], color = 'royalblue' , width = .4)
ax.bar(X + 0.4, ridge_coef['Coefficient'], color = 'purple' , width = .4)
plt.legend(labels = ['Lasso','Ridge'])
plt.show()
```

Difference between coefficients from ridge and lasso regression





In []:

In []:

In []: